Facial Detection and the Effect of Ageing

Harshita Sengar
Computer Science and Engineering
Galgotias University
Greater Noida, India
harshitasengar7597@gmail.com

Garima Pandey
Computer Science and Engineering
Galgotias University
Greater Noida, India
garima.pandey@galgotiasuniver
sity.edu.in

Abstract*—In computer vision, generating images of faces that appear to age is a challenging challenge. This is known as a conditional image generation problem because the generated image is dependent on a specified face image. Until recently, generative models were incapable of producing realistic, high-resolution photo simulations. Nevertheless, a generative model known as a generative adversarial network (GAN) has demonstrated remarkable abilities both in absolute and constrained contexts for generating realistic images. Despite this, generating images that display the aging process while preserving The true identity of the individual in the image source is a difficult task due to two contradictory constraints: The image that is produced must maintain the person's identification from the primary image. while also possessing all the characteristics of age targeted. Regarding this study, the authors used a GAN in a conditional setting in conjunction with a custom loss function designed to strike a balance between preserving the individual's identity and generating characteristics of the target age.

Facial-Aging, GAN(General Adverse Network), The CNN algorithm, Generational-Prototype, and Facial-Synthesis are index terms.

I. Introduction

Yes, facial aging is an organic process that gradually modifies an individual's appearance and it can be challenging to recognise someone at a later age based solely on their earlier photographs. Therefore, it is necessary to develop techniques capable of generating the facial image of a person at various ages, including reverse aging.

In this article, the authors present a novel method for generating facial images of a given person across six age categories.They provide examples of images generated for a particular individual at various ages. This technique has many practical applications, especially in forensic applications, such as identifying suspects or victims by enhancing or altering their photographs. Historically, this task was completed manually with the assistance of artists, computer graphics, and other techniques. Nevertheless, this procedure can be streamlined by employing an algorithm that can generate age-progressed images. Recently, The generated images have degraded in quality. The introduction of generative adversarial networks (GANs) has demonstrated remarkable abilities to generate realistic images under both constrained and unconstrained conditions. Two networks comprise GANs: A tool for generation and a discriminator(D). This is the generator network instructed to produce images that persuade the discriminator network that they are authentic. The discriminator network is taught to differentiate between actual and artificial images. By training these two networks jointly, GANs can produce images that are remarkably similar to actual images.

In the case of face aging, GANs can be utilized in a conditional setting in which the generator is trained to produce an image of the same individual at a different age based on an input image. However, the image must adhere to two constraints: It must maintain the true nature of the individual in the input picture and possess all the characteristics of the desired age. During the whole training procedure, custom loss functions can be used to satisfy these constraints.

Face degeneration is a crucial and convoluted digital imaging technology, with applications in forensic science, entertainment, and other disciplines. Recent advancements in deep learning and GANs have generated natural-looking images of a person at various ages, but there is still much work to be done to improve the accuracy and realism of these images.

This is an excellent summary of the research on face aging and how generative adversarial networks (GANs) have made it possible to generate more realistic images. The cGAN and encoder-decoder architectures have also improved the accuracy and control of the images produced. The method proposed in this study aims to generate facial images at various specified ages while maintaining the subject's identity. Using the custom loss function and GAN loss to satisfy these two constraints, the experiment demonstrates an improvement in the generation of images that preserve the source person's identity.

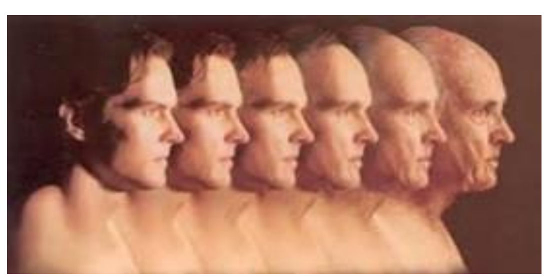

Figure 1: proposed model generates photographs of a person at distinct ages. The first column contains the user-supplied image, while the remaining columns contain images generated for each age group.

Construction:

Both the discriminator and generator models' layer descriptions are detailed in Tables 1 and 2. It corresponds to the StarGAN implementation's generator and the discriminator levels[13]. Across all layers besides the output layer, instance normalization has been implemented. As the discriminator, leaky ReLU was used. In the second table, it represents the number of domains, nc represents the dimension of the domain label, Conv represents convolution with parameters (output channels, kernel size, stride, and padding), and InstNorm represents the instance of the normalization structure L. ReLU.

Layer	Input size → Output size	Layer type
Hidden Layers	(r,c,3) → ($\frac{r}{2}$,$\frac{c}{2}$,64)	Conv(64, (4×4), 2, 1), L. ReLU
	($\frac{r}{2}$,$\frac{c}{2}$,64) → ($\frac{r}{4}$,$\frac{c}{4}$,128)	Conv(128, (4×4), 2, 1), L. ReLU
	($\frac{r}{4}$,$\frac{c}{4}$,128) → ($\frac{r}{8}$,$\frac{c}{8}$,256)	Conv(256, (4×4), 2, 1), L. ReLU
	($\frac{r}{8}$,$\frac{c}{8}$,256) → ($\frac{r}{16}$,$\frac{c}{16}$,512)	Conv(512, (4×4), 2, 1), L. ReLU
	($\frac{r}{16}$,$\frac{c}{16}$,512) → ($\frac{r}{32}$,$\frac{c}{32}$,1024)	Conv(1024, (4×4), 2, 1), L. ReLU
	($\frac{r}{32}$,$\frac{c}{32}$,1024) → ($\frac{r}{64}$,$\frac{c}{64}$,2048)	Conv(2048, (4×4), 2, 1), L. ReLU
Outout Layer (D_{adv})	($\frac{r}{64}$,$\frac{c}{64}$,2048) → ($\frac{r}{64}$,$\frac{c}{64}$,1)	Conv(1, (3×3), 1, 1)
Outout Layer (D_{cls})	($\frac{r}{64}$,$\frac{c}{64}$,2048) → (1,1,n_d)	Conv(n_d, ($\frac{r}{64} \times \frac{c}{64}$), 1, 0)

Table 1. Construction of Generator

Layer	Input size → Output size	Layer type
Down-sampling	(r,c,3+n_c) → (r,c,64)	Conv(64, (7×7), 1, 3), InstNorm, ReLU
	(r,c,64) → ($\frac{r}{2}$,$\frac{c}{2}$,128)	Conv(128, (4×4), 1, 3), InstNorm, ReLU
	($\frac{r}{2}$,$\frac{c}{2}$,128) → ($\frac{r}{4}$,$\frac{c}{4}$,256)	Conv(256, (4×4), 1, 3), InstNorm, ReLU
Bottlenack	($\frac{r}{4}$,$\frac{c}{4}$,256) → ($\frac{r}{4}$,$\frac{c}{4}$,256)	Residual Connection: Conv(256, (3×3), 1, 1), InstNorm, ReLU
	($\frac{r}{4}$,$\frac{c}{4}$,256) → ($\frac{r}{4}$,$\frac{c}{4}$,256)	Residual Connection: Conv(256, (3×3), 1, 1), InstNorm, ReLU
	($\frac{r}{4}$,$\frac{c}{4}$,256) → ($\frac{r}{4}$,$\frac{c}{4}$,256)	Residual Connection: Conv(256, (3×3), 1, 1), InstNorm, ReLU
	($\frac{r}{4}$,$\frac{c}{4}$,256) → ($\frac{r}{4}$,$\frac{c}{4}$,256)	Residual Connection: Conv(256, (3×3), 1, 1), InstNorm, ReLU
	($\frac{r}{4}$,$\frac{c}{4}$,256) → ($\frac{r}{4}$,$\frac{c}{4}$,256)	Residual Connection: Conv(256, (3×3), 1, 1), InstNorm, ReLU
	($\frac{r}{4}$,$\frac{c}{4}$,256) → ($\frac{r}{4}$,$\frac{c}{4}$,256)	Residual Connection: Conv(256, (3×3), 1, 1), InstNorm, ReLU
Up-sampling	($\frac{r}{4}$,$\frac{c}{4}$,256) → ($\frac{r}{2}$,$\frac{c}{2}$,128)	DeConv(128, (4×4), 2, 1), InstNorm, ReLU
	($\frac{r}{2}$,$\frac{c}{2}$,128) → (r,c,64)	DeConv(64, (4×4), 2, 1), InstNorm, ReLU
	(r,c,64) → (r,c,3)	Conv(3, (7×7), 1, 3), Tanh

Table 2. Construction of Discriminator

II. Related Work

To assess whether this disparity within facial recognition capability among young and elderly participants is influenced by the age of the face, an experiment was conducted in which young as well as elderly participants studied and identified faces of differing ages. The findings revealed that age-associated deficiencies within accurate recognition diminished for older features, and this effect extended to face-picture recognition. However, face age had no effect on the aging-related rise in incorrect face recognitions. This indicates that the ability to accurately recognise features and the propensity to generate false alarms are separate processes that are affected differently by aging.

This article from [2] discusses the phenomenon of cognitive aging and emphasizes several significant findings that must be considered To understand the reasons for age-related variances in cognitive performance. Contrary to popular belief, the article notes that age-related cognitive declines are relatively significant and commence in early adulthood. In addition, these declines are observed in a variety of cognitive abilities and are not always accompanied by an increase in inter-individual variability. These results indicate that cognitive aging is a complex and multifaceted process that cannot be solely attributed to age-related brain alterations. Alterations in motivation, experience, and environment may instead play a role in shaping cognitive distinctions associated with aging. Overall, the article emphasizes the need for additional research to better comprehend the mechanisms underlying cognitive aging and to develop interventions that can mitigate its negative effects.

Using deep autoencoder networks, the paper [3] proposed a method to convert high-dimensional data to low-dimensional codes. To rebuild trajectories with significant dimensionality information, autoencoder networks are trained with A neural network that consists of multiple layers containing A light internal wrapping. However, the network's efficacy is contingent on the initial weights. With the goal to resolve the same issue, this paper describes a productive approach for initializing Weights are factors that allow deep autoencoder networks to learn low-dimensional codes more efficiently compared to the key elements of research. at reducing the dimensionality of data. This method is particularly advantageous while encountering data that is highly dimensional, for instance astronomical observations images, that must be represented in a condensed and meaningful manner. Experiments on multiple benchmark datasets demonstrate which the suggested approach surpasses existing methods with respect as well as accuracy and efficiency. Overall, this work presents a practical solution for dimensionality reduction in high-dimensional data and makes a significant contribution to the field of deep learning.

Combining autonomously-organizing maps and CNNs are convolution neural networks, the authors of [4] proposed a hybrid neural-network approach for human face recognition. The SOM reduces all dimensionality for image samples and provides invariance against minor changes, whereas the CNN supplies some resistance to rotation, translation, expanding, and distortion. In a series of stratified strata, CNN is designed to derive increasingly large features. Utilizing the Karhunen-Loeve evolved as opposed to autonomously-organizing maps and aMLP (multilayer perceptron) as opposed to the CNN, a particular system is compared. The results are demonstrated using the database containing four hundred pictures of forty personalities with varying expressions, poses, and facial characteristics. The authors analyze The difficulty of computation as well as how the trained recognizer may have been developed and include new classes and systems described in this paper compared favorably to other face recognition methods.

The purpose of [5] was to investigate the reliability of eyewitness identification in a simulated criminal scenario. Three groups of eyewitnesses observed two targets and were subsequently assessed for identification accuracy 2, 21, or 56 days later. A live "show-up" or photographs were used to assess the subjects, and only one of the two objectives were present within the five-man group during testing. According to [5], there was a correlation between delay and the number of false alarms, with lengthier delays resulting in more false alarms. The testing method affected the number of matches, with live appearances yielding more accurate identifications than photographs. More than sixty percent of individuals chose among the 4 distractor individuals, indicating a high rate of misidentification. Only 28% of the subjects made no identification mistakes. These findings in [5] have significant implications for criminological application of witness testimony, as they suggest that eyewitnesses may make erroneous identifications even with relatively brief delays. Additionally, the study highlights the potential limitations of using photographs to identify suspects.

III. Proposed System

Python was used to implement the model using PyTorch [4]. NVIDIA 1060 GPU was used for all investigations.

A. Dataset

Age-labeled facial images are necessary for training the model. UtkFace [5] dataset was utilized for this purpose. This dataset includes over 23,000 pictures of Humans with ages varying between 0 and 116. These pictures were separated into six age categories. 0-18, 19-30, 31-40, 41-50, 51-60, and 60+. After separating the data into age categories, there are unequal quantities of data in each age category. In order to accomplish the above, they physically supplemented data by limiting the facial region utilizing an expression analyzer.

B. Laboratory research Set-up

Model input image measurements are 128 by 128 pixels. The model was trained using the Adam optimisation technique[6] opposed to GDO. It has been shown to be the best option for working with GANs. Beta1 and beta2 are two factors of the Adam optimizer(AO) that determine the gradient descent's momentum. We built the model with a group size of eight for fifty iterations. Few additional hyper-parameters for the model are listed by the method.

C. Evaluation Matrix

In any context, the generated picture is difficult to evaluate. No widely acknowledged method exists for assessing the fidelity of pictures produced. GAN is a model currently in utilization to generate various picture forms during recent years. Various publications had expressed results in various methods. The majority are human evaluations [10]. A volunteer is shown both a computer-generated and an actual image for a brief period in time. The goal is to determine which is artificial and which is genuine? This grade has been converted to reflect the image's performance. This evaluation lacks both credibility and comparability.

For the purpose of protecting the personal information of the original individual, we created images of various age categories. FaceAging [7] had previously introduced a Face detection value evaluation matrix used to evaluate the task of maintaining the first someone's name. An already trained algorithm for visage confirmation is applied for this purpose. OpenFace [8] has been employed for face authentication. The OpenFace platform is a freely downloadable library capable of recognising or validating features.

The face authentication algorithm(FVA) admits two images as input and gives a score that denotes the degree of match between the two identical faces. The reduced score, greater is its confidence that both images show exactly the same individual. In general, a separation of 0.70 or lower is sufficient over beneficial evidence. We determined the distance between all generated images of varying ages and the original images. Because every picture in which this distance was below 0.70, we declared it a positive test and computed the accuracy of the entire model.

Due to the fact we've produced images of multiple ages, the image needs to maintain its features of the corresponding category of age. For this evaluation, a classifier of these age categories has been trained. Based on the generated image and the aggregate accuracy of each class, we computed the predicted class for each image.

D. Face Recognition Score

Using the pre-trained OpenFace library, the visage identification rank was calculated, similarly described within the preceding segment. We already compared their outcome to known prior outcomes. Antipov et al. [7] proposed a face aging method that employs explicit Identity preserving optimization to preserve the individual's identity. They failed to prioritize the rate of categorization or the quantity of pictures with material.

IV. Tools and Technology

There are a variety of available tools and technologies for facial detection and aging effects. Popular ones include:

OpenCV: OpenCV is an open-source, free artificial intelligence and image processing algorithm. Face detection, face recognition, and age estimation modules are pre-installed.

DLib: DLib is an artificial intelligence and image processing C++ library, which includes pre-built models for face recognition, face landmark recognition, and age estimation.

Tensorflow: Google developed the open-source machine learning framework Tensorflow. Face detection, face recognition, and age estimation are pre-built models.

PyTorch: Facebook developed the open-source machine learning framework PyTorch. Face detection, face recognition, and age estimation are pre-built models. FaceApp is a mobile application that applies various facial effects, including aging effects, to photographs using neural networks. AgingBooth is a mobile application that simulates the aging process on a person's visage using face detection and aging algorithms.

V. Methodology

The basic idea of the generating artificial network, a key component of our model, gets clarified first.. Finally, we present our future generation's approach.the visage images of individuals of various ages. The impact of the numerous objective function components is then discussed at last to explain the model's training processes.

A. Network of Creative Writing Challengers

In attempting to trick a discriminator, the programmer learns to produce pictures that approach the target domain, which has been trained to differentiate between genuine and synthetic images. Training is provided to the generating device and detector iteratively using adversarial training, where the generator attempts to improve its performance by generating better images and the discriminator attempts to improve its ability to distinguish between actual and generated images. This procedure is repeated until the generator is able to produce images that the discriminator cannot distinguish from actual images.
GAN wants to create an actual, unique image that matches the distribution of the domain under all circumstances.

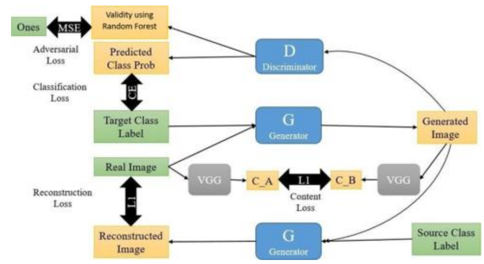

Fig. 2 Random Forest Forest Training

The illustration in Figure 2 describes GAN development. GAN can also be utilized in conditional contexts by providing the generator with additional input. The constrained generated artificial network (cGAN) was introduced by Mirza et al. in [9]. By providing additional information, such as the target class or attribute, GAN can be utilized in the conditional configuration.
It popularized the image-to-image translation dimension of the generative model.These models generate a target image using an input image as a condition. Isola et al. [10] first proposed the network for image-to-image conversion which generates an excellent quality photograph. Because an image is used as a conditional variable, This type of model is generated using encoder-decoder architecture.

The compression algorithm receives the image as contributions and transforms it into a matrix with fewer parameters, which is then used as input by the receiver. It is sampled by the receiver in order to generate the desired image. Occasionally, the algorithm for decoding accepts an unpredictable vector as input in addition to the encoding device matrix. The purpose of the unpredictable vector is to imbue the generated image with inconsistencies.

B. Aging Face prototype

Our goal is to educate a model for picture-to-picture transformation across multiple domains. In addition, the contentious and customary losses were implemented to make the image resemble the target domain while preserving the identity of the source. The discriminator functions as a supplementary classifier [2], enabling a single discriminator to evaluate images from multiple fields. A discriminator has two outputs to this objective: class probability and software validation. Similar to traditional GAN, validity is the random score of an image living authentic or fraudulent, category probability is the expected value of a derived image being a part to a particular area, as suggested using a discriminator.

Domains Classifier Changes: The objective of this framework is to offer a picture of the area of focus that is recognised as a representation associated with the domain.. We already have implemented For this, the discrimination algorithm produces an additional result and an assessment loss. The classifier is given instructions on how to correctly identify pictures during distinction training.

Reconstruction Loss: By minimizing this hostile loss teaches the generator how to create a realistic image. that corresponds to the target image's statistics division. The classification reduction creates visuals which resemble those of the subject class. However, none of these losses impose any restrictions on keeping the source image's details or preserving the source image's subject's identity. To add this constraint, we optimize two additional generator losses.

The reconstruction loss imposes half of CycleGAN's objective [11]. Due to the fact that the curriculum level is a parameter to a generator, the two separate ones are unnecessary to impose this loss.This loss provides an indirect requirement on the image that is produced so it can be transformed back to the initial picture with the identical class level. Despite the fact that this loss demonstrated a remarkable result according to reports within an assortment of uses, of papers, it was a significant setback. Nonetheless, when the engine discovers an easy way that produces the picture returned from its starting nation, known as picture might get converted with no the feature you want being preserved at the target level. We have implemented content loss to guarantee that the material of the original image is maintained in the image that was generated.

Content Loss. The intent of content loss is to require the generator to preserve the fundamental content of the source image in the generated image. It is the difference between the content of the source file and the generated image. comparable losses in comparable programmes, such as image super-resolution, have yielded positive results [4]. Extracting the data information via a deep convolutional neural network's learnt high-level representation. Deep CNNs trained on huge data sets are well-known to be efficient at identifying high-level image features. In order to detect content loss, we used the Image Network [6]-trained visual Geometric Group-19 [5] model for obtaining characteristics.

The generator's output is provided with target identifiers for subclass and real image. The picture that is generated is then sent along with the source class label to the detector and maker. Validity and predicted class probability are output by the discriminator. We compute two objective losses based on this information: both the conflict and classification losses. The algorithm that generates it results in the unprocessed image's rebuild. It is to determine the difference involving the initial pic and the picture that was reconstructed with L1 loss.The real picture and its associated generated pictures are fed into the trained fully stored visual geometry group-19 algorithm to determine information degradation. What kind of damage occurred is then determined by combining the results of both models.

C. Construction of Networks

Author in [7] presented cGAN after the introduction of generative adversarial network for low-resolution images. This model's architecture was adopted, with modifications and variations incorporated into subsequent publications. [8] Both the generator and discriminator utilize a convolution group to renormalize the ReLU module.

D. Coaching: In order to sustain the GAN training and produce high-quality images, a modified training objective was proposed. Wasserstein [2, 3] demonstrated that optimizing the GAN objective proposed in the original GAN paper [7] is challenging. The proposed second WGAN is known as WGAN-GP. As an alternative to clipping, gradient plenty was utilized to enforce Lipchitz constant. A differentiable function f is 1-Lipschitz if the maximal value of its gradient everywhere is ohe..

Consequently, instead of clipping the model as WGAN does, WGAN- GP penalizes it ifThe gradient norm is not identical to its objective norm of one. In the course of our model's instruction.

Generator with bypass link: The issue of excellent quality photographs are used in the process of translating a photograph into another photograph. In the majority of Both the source & target pictures have a comparable visual format when used in solutions despite their distinct what it looks like. In the app that we are developing, the identity of a person who differs at various points of ages must be preserved, which is why we developed the loss function. We believe that this property is the generator.

Numerous applications in this field employ encoder-decoder framework [9]. In this kind of network, the input is transmitted to the generator so that it can be compressed to create a low-dimensional approximation. This concept is referred to as block visualization.. The data is further transmitted to the decoder, which resamples it to the image size. In numerous applications, the majority Both networks share some fundamental characteristics. Therefore, it would be beneficial to instantaneously transmit the same attribute between networks.

We adhered to the logic presented in U-Net [2] in order to provide the source who has the tools to spread this knowledge. In between encoder and decoder layers, this network utilizes skip-connection. ResNet [2] popularized the idea of skipping the connection. It utilized a connection among its tiers that skips. Combining inputs of a layer with its output or the outputs of multiple layers is the concept. It transmits direct information between network layers to aid the model in understanding a transition's transition and change. A bypass level i to level (n-i), wherein n represents the total amount of levels within the a generator, are connected, exists in U-Net, is inserted between the encoder and decoder.

VI. Conclusion

TNumerous applications exist for the ability to generate visage images of various ages based on a picture and the present age-factor. In this paper, the method for completing this endeavor is proposed. Our GANs-dependent procedure is able to produce facial pictures of different peer-groups. Experiments have proven the fact that our approach is superior and preserves the individual's identity in the generated image than other methods. In addition, our method produces images with age-appropriate characteristics, which is supported by the precision of the categorization. Suggested enhancements for this effort involve incorporating photos via greater pixels to enhance the picture quality. and enhancing the image recognition function.

VII. Result and Discussion

This research contributes to the field of facial expression recognition by examining aging's effect on how one looks at detection. Our experiments demonstrate that age-related effects must be considered when attempting to recognise facial expressions. By considering this, the solution for expression recognition will be more general and efficient.

Fig. 6. Illustrations of pictures produced by the suggested technique.

The programme is written in Python and its graphical interface is provided by the Tkinter module. It uses the NLTK for text processing, as well as Math and PyPDF2. The programme prompts the user to choose a YouTube video to summarise, retrieves the transcription, and saves the file as a PDF. The preparatory phase involves tokenization, removal of stop words, and stemming. The algorithm then employs a statistical approach to frequency-inverse document frequency to characterize the document's characteristics and selects key words and phrases based on these characteristics. The overview is displayed in the interface segment. Using Rouge 2.0, which employs a synonym dictionary to capture semantic overlap, the efficacy of the programme is evaluated. Using the CNN/daily mail dataset, the results of the programme are compared to other traditional methods and baselines, and it is discovered that the programme significantly outperforms them with a higher score.

REFERENCES

[1] Fulton, A. and Bartlett, J.C., 1991. Young and old faces in young and old heads: the factor of age in face recognition. *Psychology and aging*, *6*(4), p.623.

[2] Salthouse, T. A. (2004). What and when of cognitive aging. Current Directions in Psychological Science, 13, 140-144.

[3] Hinton, G.E. and Salakhutdinov, R.R., 2006. Reducing the dimensionality of data with neural networks. *science*, *313*(5786), pp.504-507.

[4] Lawrence, S., Giles, C.L., Tsoi, A.C. and Back, A.D., 1997. Face recognition: A convolutional neural-network approach. *IEEE transactions on neural networks*, *8*(1), pp.98-113.

[5] Goldstein, A.G., 1977. The fallibility of the eyewitness: Psychological evidence. *Psychology in the legal process. New York: Spectrum*.

[6] Mason, S.E., 1986. Age and gender as factors in facial recognition and identification. *Experimental Aging Research*, *12*(3), pp.151-154.

[7] Johnson, M.H., Senju, A. and Tomalski, P., 2015. The two-process theory of face processing: modifications based on two decades of data from infants and adults. *Neuroscience & Biobehavioral Reviews*, *50*, pp.169-179.

[8] Suo, J., Zhu, S.C., Shan, S. and Chen, X., 2009. A compositional and dynamic model for face aging. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, *32*(3), pp.385-401.

[9] Tazoe, Y., Gohara, H., Maejima, A. and Morishima, S., 2012. Facial aging simulator considering geometry and patch-tiled texture. In *ACM SIGGRAPH 2012 Posters* (pp. 1-1).

[10] Bartlett, J.C. and Fulton, A., 1991. Familiarity and recognition of faces in old age. *Memory & Cognition*, *19*(3), pp.229-238.

[11] George, P.A. and Hole, G.J., 1995. Factors influencing the accuracy of age estimates of unfamiliar faces. *Perception*, *24*(9), pp.1059-1073.

[12] Agrawal, A.K. and Singh, Y.N., 2015. Evaluation of face recognition methods in unconstrained environments. *Procedia Computer Science*, *48*, pp.644-651.

[13] Fontaine, X., Achanta, R. and Süsstrunk, S., 2017, March. Face recognition in real-world images. In *2017 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)* (pp. 1482-1486). Ieee.

[14] Puthea, K., Hartanto, R. and Hidayat, R., 2017, November. A review paper on attendance marking system based on face recognition. In *2017 2nd International conferences on Information Technology, Information Systems and Electrical Engineering (ICITISEE)* (pp. 304-309). IEEE.

[15] Durkin, E., 2019. New York school district's facial recognition system sparks privacy fears. *The Guardian*.